

HTA Firmware Tracker


A web dashboard that tracks the latest firmware/software versions for HTA's audio gear, so nobody has to manually check manufacturer sites.

What it does

- Tracks 47 pieces of equipment across DiGiCo, Yamaha, Solid State Logic, Allen & Heath, Shure, d&b audiotechnik, and Dante/Audinate.
- **All 47 items are checked automatically** — most by scraping a static page, several manufacturers via a real headless browser where the site needs JavaScript or blocks plain requests (see [How each manufacturer is checked](#)).
- Re-checks everything once a day (via a scheduled GitHub Actions workflow - see [Deployment](#) - or via a background job if you run it locally) and highlights anything whose version changed since the last check.
- A stats panel at the top (Tracked / Updates / Errors / Last checked, the last shown as a relative time like "2h ago"), with a search box and Up to date / Update available / Check failed filters. Clicking the Updates card jumps straight to the flagged rows.
- Dashboard grouped by manufacturer, with status badges (Up to date / Update available / Check failed). A row with a new version gets a soft amber highlight in place (it isn't reordered) that clears itself automatically 1 month after the change was first detected — no one needs to click anything to dismiss it. Installable as a PWA on a phone (Add to Home Screen) for a full-screen, app-like view.
- Items that are genuinely the same underlying platform tracked from one source — e.g. DiGiCo's six Quantum consoles — collapse into a single family row ("Quantum series", under DiGiCo's own section heading) instead of listing near-duplicate rows that always move together. See [Firmware families](#) for how that grouping is decided and kept in sync with Slack notifications.

Architecture at a glance

One flow, run in two places (a daily GitHub Actions job, and locally on demand):



- **seed.py is the list of what's tracked;** each row names a `checker_key` (e.g. `digico:quantum`), whose prefix `registry.py` routes to a manufacturer module. A checker gets all of its items at once and returns one `CheckResult` per item - it fetches each *source* once, not once per model.
- **State lives in git, not a database server.** CI starts from nothing each run: it re-seeds a throwaway SQLite DB, overlays the last known versions/timestamps from `data/equipment.json` (matched on manufacturer + model), checks, and writes the result back. So history persists as commits, and "what changed" is a before/after diff of two JSON snapshots (`app/changes.py`) rather than anything stored.
- **The deployed site is static HTML.** `build_static.py` renders the same Jinja template and calls the same `build_dashboard_context()` as the live FastAPI app, so the two can't drift on grouping/sorting/stats. There's no server in production - and so nothing that writes: flags clear themselves after a month instead of needing an "acknowledge".
- **Failures degrade, they don't cascade.** A checker that raises marks only its own items "Check failed"; a Slack outage or missing webhook never blocks the daily commit; a release-date lookup failing doesn't fail the version check it rode along with.

Quick start

```

cd "hta-firmware-tracker"
python3 -m venv .venv
source .venv/bin/activate
pip install -r requirements.txt
playwright install chromium
uvicorn app.main:app --reload

```

Open `http://127.0.0.1:8000` — the database (SQLite, at `data/tracker.db`) is created and seeded automatically on first run, and a check kicks off ~5 seconds after startup. (The `playwright install chromium` step downloads a ~90MB headless browser used only by the Allen & Heath / Shure / d&b / Dante checkers — see below.)

How each manufacturer is checked

Manufacturer	Method	Notes
DiGiCo	✓ Scraped	<code>support.digico.biz</code> 's Zendesk help-center JSON API
Yamaha	✓ Scraped	Each product's Downloads page on <code>usa.yamaha.com</code> (a <code>#firmware-table</code> with a genuine "Last Update" date per file, not just a version)
Solid State Logic	✓ Scraped	<code>support.solidstatelogic.com</code> 's Zendesk API too; tracked via SOLSA (versions 1:1 with SSL Live console software) and the Network I/O firmware bundle
Allen & Heath	✓ Scraped (headless)	Version: <code>www.allen-heath.com</code> returns HTTP 403 to plain requests (bot protection) — works fine in a real headless browser. Release date: a separate evergreen article on <code>support.allen-heath.com</code> (a Zendesk help centre, read via its public JSON API — even a real browser can't get past that subdomain's own Cloudflare challenge page), month/year only, no day
Dante/Audinate	✓ Scraped (headless)	The version table is inside a click-to-expand accordion, so it's not in the page until JS runs. Dante Controller is published as separate Windows / macOS Apple Silicon / macOS Intel builds with their own versions, so all three are tracked (see below)
d&b audiotechnik	✓ Scraped (headless)	Their Download Center is a JS-driven search box behind a cookie-consent overlay; D40/D90 share one firmware release, DN1 Switch has its own
Shure	✓ Scraped (headless)	Firmware versions come from <code>shure.com</code> 's searchable software/firmware archive listing

Every item is checked automatically. A piece of gear only gets added to the list once a checker can read its version.

Per-platform versions (Dante Controller). Its row stays compact - one status badge and the newest version, with a "3 platforms" note - and expands (chevron on desktop, tap on phone) to list each platform's own version. Each platform tracks its own previous version and change date, and the row's badge is "Update available" whenever *any* platform changed within the 1-month highlight window (the expanded list tags which one). Slack alerts name the platform(s), e.g. "Dante Controller (Windows) → 4.18.1.3", and an update to just one platform is caught even if it doesn't change the headline version. Each platform also has its own real release date - Audinate publishes a dedicated "Release Notes" page per platform build, linked right from the same accordion the version comes from - so the expanded list shows each platform's own date, and the row's own Release date column shows whichever platform's date matches the headline version.

Article-based sources (DiGiCo, SSL). These manufacturers announce firmware in a help-centre *article* rather than on a dedicated firmware page, so a new release shows up as a new article/title/version and is handled like any other update - same title/version-based detection as every other checker. The DMI-Dante64@96 card is tracked as two rows (Zynq HC and the older Summit HC), since they're separate hardware variants with their own firmware lines - one shared row used to flip between whichever line published most recently.

Every manufacturer publishes a real release date now. The "Detected on" fallback (the day the tracker itself first caught a version change, `last_changed_at`, not a manufacturer-published date) is still there for whenever a future manufacturer's checker has nothing to put in `release_date` - it's determined per manufacturer from the data (any item with a real release date), not a hardcoded list, so nothing needs to change here if that happens again.

How the headless-browser checkers work

`app/checkers/browser_base.py` holds a small Playwright helper (`browser_page()` launches a headless Chromium tab, `dismiss_cookie_banner()` clicks past consent overlays). `allenheath.py`, `dante.py`, `dbaudio.py`, and `shure.py` use it to load a page, interact with it if needed (click an accordion, type into a search box), and pull the version out of the rendered text. They're slower than the plain-HTTP checkers (each spins up a real browser) but run fine as part of the once-a-day background check. To add another one, copy the shape of `dante.py`, then register it in `app/checkers/registry.py` and point the relevant `app/seed.py` rows at it with `"scrape"` + a `checker_key`.

Firmware families

Some rows track the same underlying platform from one firmware source - DiGiCo's six Quantum consoles, Allen & Heath's four dLive surfaces/ MixRacks, Yamaha's two Rio-D2 I/O racks, Yamaha's eight Rivage PM components, d&b's three D-series amps - and always report the identical version because they're assigned the same `CheckResult` from one `check_all()` call. Those collapse into a single dashboard row (e.g. "Quantum series", under DiGiCo's own section heading) instead of six near-identical ones. On desktop/tablet every field - status, version, release date, previous version, last checked, source - still shows directly in that one row, since it's identical for every member; only the individual model names (e.g. Quantum 112/225/326/338/5/7) sit behind the chevron, as pill tags. On phone, where the whole row's detail is already a tap away, the model pills join the rest of that existing reveal instead of getting their own control. A family whose collective name already spells out every member - d&b's amp group is named "D25/D40/D90", not a generic "series" label, because unlike the others it has no shorter collective name that wouldn't also fit some other d&b product - skips the chevron/expand entirely on both desktop and phone, in `FAMILY_NO_EXPAND` (`app/families.py`): there's nothing left for it to reveal.

This is deliberately **not** based on matching version-number strings - two unrelated products could coincidentally share a version number (Shure's AD610, SBC240, and ADXR are different product types and always stay as separate rows even if that happens). The real signal is `checker_key` (assigned per actual firmware source in `app/seed.py`): items sharing one were checked together and are guaranteed to move together. `app/families.py` holds the one `FAMILY_NAMES` mapping (`checker_key` → friendly collective name) that both `app/dashboard_data.py` (row consolidation) and `app/slack.py` (notification grouping, see [Slack notifications](#)) import, so the two can never disagree about what counts as "the same family." Those names are deliberately brand-free ("Quantum series", not "DiGiCo Quantum series") - on the dashboard the row already sits under its manufacturer's own section heading, so repeating the brand on every row under it would be redundant. Slack messages have no such heading, so `app/slack.py` prepends the manufacturer back on when it builds the notification text (correcting Yamaha's all-caps `YAMAHA` field to "Yamaha" through

`MANUFACTURER_LABELS` along the way) - a Slack message still reads "DiGiCo Quantum series" even though the stored name is just "Quantum series".

A family whose members all share one dashboard category (Yamaha's two Rio-D2 racks are both "I/O Racks") is filed under it normally, same as any other row. One spanning several categories (Yamaha's Rivage PM components live under Consoles, DSP Engines, and I/O Racks) isn't attributed to any one of them arbitrarily - it instead gets its own dedicated subsection, named in `FAMILY_CATEGORIES` (also in `app/families.py`), e.g. "Rivage PM Series". That's only set for families where it's actually useful - one that's the only thing tracked under its manufacturer (Allen & Heath's dLive series) stays without a heading instead, since a heading identical to the one row underneath it would just be noise. A manufacturer's subsections don't have to sort alphabetically or by equipment type either: Yamaha's "Rivage PM Series" is pinned ahead of "Consoles" in `CATEGORY_ORDER` (`app/dashboard_data.py`) since it's their flagship line, not because any general ordering rule would put it first.

DiGiCo's SD10 is a genuine singleton (no other SD-series console is tracked), but it still renders as a family-style row - "SD series", with a chevron expanding to a lone "SD10" pill - rather than falling back to a plain, non-expandable row the way every other singleton does. That's `FAMILY_FORCE` in `app/families.py`: a `checker_key` can opt into the family treatment even with one current member, when (like DiGiCo SD) its `FAMILY_NAMES` entry represents a whole product line rather than describing one specific model. Its dashboard category is just "Consoles" though, same as the Quantum family right above it - no dedicated subsection, since both are genuinely the same equipment type.

A family's members are supposed to always report identically (same `CheckResult`, one `check_all()` call), but that premise could still be violated by something outside any single check run - a member added to the family before its first real check, stale state from before it joined, a future checker bug. So `_consolidate_families` doesn't just trust the first member for the row's displayed status: it picks whichever member's status is most urgent (error beats a pending update beats "ok"), so a failing member can't be quietly absorbed into a row that otherwise looks fine. A family with any failing member shows "Check failed" on its row. The error line underneath is the message once when the whole family failed together (the normal case - one shared fetch), and names specific models only when just some failed or the messages differ. There's deliberately no per-model marker in the expanded pills - with every current checker assigning one shared result per `checker_key`, a single member failing alone can't happen in practice, so the row-level status is all that's needed.

The stats panel's two counts deliberately mean different things. "Tracked" is a coverage number - every real, individually-tracked piece of equipment, unaffected by how many rows that collapses into on screen, so it doesn't shrink just because some of what it covers shares one firmware source. "Updates" counts rows instead, not units - if Rivage PM's 8 components all pick up a new version at once, that's genuinely one thing to go look at, so it shows "1 update available," not 8. Clicking the Updates card filters straight to those rows, so the number it shows always matches what you'd actually see land on screen. "Errors" still counts units (unchanged) - see `build_dashboard_context` in `app/dashboard_data.py` for exactly which list each one sums over.

Deployment

The published site (on Vercel) is a **static build**, not the live FastAPI app — there's no server running continuously anywhere. Instead:

- `.github/workflows/daily-check.yml` runs once a day (GitHub Actions cron, plus a manual "Run workflow" button). It rebuilds a throwaway database from `data/equipment.json`, runs every checker for real, writes the results back to `data/equipment.json`, renders the dashboard to `public/index.html` (`scripts/build_static.py`), and commits both back to the repo — a fresh "Last checked" timestamp lands every single day, even when no version actually changed, so there's always a commit to show the check ran.
- **Vercel** serves the `public/` directory as-is (see `vercel.json`) — no build step, no server, just static files. Importing this repo into Vercel needs no extra configuration.
- The dashboard has no interactive controls that write anything (no manual "Log" or "Acknowledge" buttons) — it's a pure status *display*, matching what a backend-less static site can actually support. Update flags clear themselves automatically instead (see below).
- `data/tracker.db` (SQLite) is still what the local interactive app uses day-to-day and is gitignored, same as always - it's not part of the deployment at all.

Slack notifications

The daily workflow posts to Slack (`app/slack.py`) whenever it finds a **genuinely new** firmware version - comparing each item's version before and after that day's run, not just whatever's currently flagged. A version that's still pending from a previous day, or one that just auto-cleared after its 1-month window, doesn't trigger a repeat message; only an actual change does.

Items are grouped by family before posting - e.g. all 6 DiGiCo Quantum consoles updating together becomes one message ("New firmware available for *DiGiCo Quantum series* → V23"), not six. Grouped by `checker_key` (items that share one get checked together and always report the same new version, by construction - a more reliable "these are really the same family" signal than just matching version-number strings, which could coincidentally collide between two unrelated manufacturers), using the same `FAMILY_NAMES` mapping in `app/families.py` that the dashboard uses to consolidate rows - see [Firmware families](#) - so the two never disagree about what counts as one family. A genuinely unrelated update detected the same day (different `checker_key`) always gets its own separate message, `@channel` and all - never bundled into someone else's.

Reads the webhook URL from the `SLACK_WEBHOOK_URL` repository secret (GitHub → Settings → Secrets and variables → Actions), passed to the workflow step as an env var - never hardcoded, and not logged anywhere (GitHub also masks any exact occurrence of a registered secret's value in the Action's log output automatically). Missing the secret, or the POST itself failing, never breaks the daily commit - `notify_updates()` in `app/slack.py` degrades to a no-op (logged, not raised) in either case.

Local automatic checking

If you run this locally instead of relying on the deployed site, the background scheduler (`app/scheduler.py`, `APScheduler`) only runs **while the app's process is alive** — if it's just open on your laptop, it only checks

while your laptop has it open. The deployed site doesn't have this limitation, since GitHub Actions runs the check regardless of whether anyone's computer is on.

When a version changes, the item gets an "Update available" badge and an amber row highlight, in its normal position - it isn't moved to the top. Both clear themselves automatically 1 month after the change was *first* detected (`Equipment.last_changed_at` , see `app/runner.py` 's `UPDATE_HIGHLIGHT_WINDOW`) - re-confirming the same pending version on later daily checks doesn't restart that clock, only a genuinely new version change does. No manual acknowledgment involved anywhere.

Project layout

```
app/
  main.py          FastAPI routes + dashboard
  dashboard_data.py Grouping/sorting/summary logic shared by the live app
                    and the static-site builder, incl. family row
                    consolidation (see Firmware families)
  families.py      FAMILY_NAMES: checker_key -> brand-free family name,
                    shared by dashboard_data.py and slack.py (the latter
                    prepends the manufacturer back on)
  models.py        Equipment / CheckLog tables (SQLAlchemy)
  runner.py        Runs checkers, diffs versions, writes history,
                    auto-expires update flags after UPDATE_HIGHLIGHT_WINDOW
  scheduler.py     Background interval job (local interactive use only)
  seed.py          The 47-item equipment list + which checker covers each
  export.py        Equipment <-> data/equipment.json round-trip, used by
                    the CI pipeline
  changes.py       Before/after diff of two JSON snapshots -> the genuine
                    version changes Slack is told about
  slack.py         Webhook notifications, one message per family
  checkers/
    digico.py, yamaha.py, ssl.py      Plain-HTTP (HTML / Zendesk JSON API)
    allenheath.py, dante.py, dbaudio.py, shure.py  Headless-browser (Playwright)
    browser_base.py  Shared Playwright helper
    registry.py      Routes a checker_key to its module
  templates/, static/  Dashboard UI
scripts/
  ci_check.py        Entry point for the daily GitHub Actions run
  build_static.py    Renders the dashboard to public/ for Vercel
data/tracker.db      SQLite database (created on first run, local dev only)
data/equipment.json  Git-committed source of truth for the deployed site
```

Adding equipment

Add a row to the `EQUIPMENT` list in `app/seed.py` (manufacturer, model, category, checker_key, source URL, optional notes), then restart the app (or run `python3 -m app.seed`) — `seed()` fully syncs the database to match the list: new rows are added, changed rows are updated, and rows removed from the list are deleted.